

Zhihao Zhu

Postdoctoral Fellow, Trustworthy AI & Information Systems
Hong Kong University of Science and Technology (HKUST)
zhihaozhu@ust.hk | zzh9568.github.io

Experience

- **Postdoctoral Fellow**, HKUST, advised by Prof. Yi Yang.
- **Ph.D. in Computer Science**, University of Science and Technology of China (2025), advised by Prof. Defu Lian.
- **Bachelor's degree in Computer Science**, Fuzhou University (2020).

Research Interests

Trustworthy AI; AI privacy and auditing; large language models; recommender systems.

Published Journal Article

- **Zhihao Zhu**, Yi Yang, Yangyang Fan, Defu Lian. Forget Me If You Can: Auditing User Data Revocation in Recommendation Systems. *Information Systems Research*, 2026.

Working Papers

- **Zhihao Zhu**, Jiale Han, Yi Yang. HoneyImage: Verifiable, Harmless, and Stealthy Dataset Ownership Verification for Image Models. *Management Information Systems Quarterly*, Major Revision.
- **Zhihao Zhu**, Hongyi Tang, Yi Yang, Ahmed Abbasi. Revealing Training Data Exposure in Vision-Language Large Models via Parameter Gradients. *Nature Communications*, Major Revision.
- **Zhihao Zhu**, Yi Yang. Governing Execution Risk in Agentic AI Systems: A Trajectory-Guided Framework for Red Teaming. *Information Systems Research*, Under Review.
- **Zhihao Zhu**, Yi Yang. RecShield: Output-Level Attribute Unlearning in Recommender Systems. *Information Systems Research*, Under Review.
- **Zhihao Zhu**, Yi Yang, Chenwang Wu, Defu Lian. GraphMSA: Stress Testing Graph Classification Services Against Model Stealing Attacks. *INFORMS Journal on Computing*, Under Review.

AI Conferences

- **Zhihao Zhu**, Yi Yang, Defu Lian. TDDBench: A Benchmark for Training Data Detection. *International Conference on Learning Representations*, 2025.
- Hongyi Tang*, **Zhihao Zhu***, Yi Yang. Identifying Pre-training Data in LLMs: A Neuron Activation-Based Detection Framework. *Conference on Empirical Methods in Natural Language Processing*, 2025. *Equal contribution.
- **Zhihao Zhu**, Chenwang Wu, Rui Fan, Defu Lian, Enhong Chen. Membership Inference Attacks against Sequential Recommender Systems. *The Web Conference*, 2023.

Service

Reviewer for *Information Systems Research* (ISR), *INFORMS Journal on Computing* (IJOC), NeurIPS, ICLR, and ACL Rolling Review.

Awards

- **National Graduate Scholarship (China)**, 2022.
- **KDD Cup 2023 “Next Product Generation” Challenge**, 2nd Place.